

# two\_layer\_net

September 27, 2019

## 1 Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

```
In [1]: # A bit of setup

import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

The neural network parameters will be stored in a dictionary (model below), where the keys are the parameter names and the values are numpy arrays. Below, we initialize toy data and a toy model that we will use to verify your implementations.

```
In [2]: # Create some toy data to check your implementations

input_size = 4
hidden_size = 10
num_classes = 3
num_inputs = 5

def init_toy_model():
    model = {}
    model['W1'] = np.linspace(-0.2, 0.6, num=input_size*hidden_size).reshape(input_size, h
```

```

model['b1'] = np.linspace(-0.3, 0.7, num=hidden_size)
model['W2'] = np.linspace(-0.4, 0.1, num=hidden_size*num_classes).reshape(hidden_size,
model['b2'] = np.linspace(-0.5, 0.9, num=num_classes)
return model

def init_toy_data():
    X = np.linspace(-0.2, 0.5, num=num_inputs*input_size).reshape(num_inputs, input_size)
    y = np.array([0, 1, 2, 2, 1])
    return X, y

model = init_toy_model()
X, y = init_toy_data()

```

## 2 Forward pass: compute scores

Open the file `cs231n/classifiers/neural_net.py` and look at the function `two_layer_net`. This function is very similar to the loss functions you have written for the Softmax exercise in HW0: It takes the data and weights and computes the class scores, the loss, and the gradients on the parameters.

Implement the first part of the forward pass which uses the weights and biases to compute the scores for all inputs.

```

In [3]: from cs231n.classifiers.neural_net import two_layer_net

scores = two_layer_net(X, model)
print(scores)
correct_scores = [[-0.5328368, 0.20031504, 0.93346689],
                  [-0.59412164, 0.15498488, 0.9040914 ],
                  [-0.67658362, 0.08978957, 0.85616275],
                  [-0.77092643, 0.01339997, 0.79772637],
                  [-0.89110401, -0.08754544, 0.71601312]]

# the difference should be very small. We get 3e-8
print('Difference between your scores and correct scores:')
print(np.sum(np.abs(scores - correct_scores)))

[[-0.5328368  0.20031504  0.93346689]
 [-0.59412164  0.15498488  0.9040914 ]
 [-0.67658362  0.08978957  0.85616275]
 [-0.77092643  0.01339997  0.79772637]
 [-0.89110401 -0.08754544  0.71601312]]
Difference between your scores and correct scores:
3.848682278081994e-08

```

## 3 Forward pass: compute loss

In the same function, implement the second part that computes the data and regularization loss.

```
In [4]: reg = 0.1
        loss, _ = two_layer_net(X, model, y, reg)
        correct_loss = 1.38191946092

        # should be very small, we get 5e-12
        print('Difference between your loss and correct loss:')
        print(np.sum(np.abs(loss - correct_loss)))
```

Difference between your loss and correct loss:  
3.605849741239453e-06

## 4 Backward pass

Implement the rest of the function. This will compute the gradient of the loss with respect to the variables  $W1$ ,  $b1$ ,  $W2$ , and  $b2$ . Now that you (hopefully!) have a correctly implemented forward pass, you can debug your backward pass using a numeric gradient check:

```
In [5]: from cs231n.gradient_check import eval_numerical_gradient

        # Use numeric gradient checking to check your implementation of the backward pass.
        # If your implementation is correct, the difference between the numeric and
        # analytic gradients should be less than 1e-8 for each of W1, W2, b1, and b2.

        loss, grads = two_layer_net(X, model, y, reg)

        # these should all be less than 1e-8 or so
        for param_name in grads:
            param_grad_num = eval_numerical_gradient(lambda W: two_layer_net(X, model, y, reg)[0],
                print('%s max relative error: %e' % (param_name, rel_error(param_grad_num, grads[param_name])))

W2 max relative error: 5.260597e-05
b2 max relative error: 6.317662e-06
W1 max relative error: 2.493902e-05
b1 max relative error: 1.449991e-03
```

## 5 Train the network

To train the network we will use SGD with Momentum. Last assignment you implemented vanilla SGD. You will now implement the momentum update and the RMSProp update. Open the file `classifier_trainer.py` and familiarize yourself with the `ClassifierTrainer` class. It performs optimization given an arbitrary cost function data, and model. By default it uses vanilla SGD, which we have already implemented for you. First, run the optimization below using Vanilla SGD:

```
In [6]: from cs231n.classifier_trainer import ClassifierTrainer
```

```

model = init_toy_model()
trainer = ClassifierTrainer()
# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient Descent (no
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.0, learning_
                                              update='sgd', sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)

print('Final loss with vanilla SGD: %f' % (loss_history[-1], ))

```

```

starting iteration 0
Final loss with vanilla SGD: 0.940683

```

Now fill in the **momentum update** in the first missing code block inside the train function, and run the same optimization as above but with the momentum update. You should see a much better result in the final obtained loss:

```

In [7]: model = init_toy_model()
        trainer = ClassifierTrainer()
        # call the trainer to optimize the loss
        # Notice that we're using sample_batches=False, so we're performing Gradient Descent (no
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                              model, two_layer_net,
                                              reg=0.001,
                                              learning_rate=1e-1, momentum=0.9, learning_
                                              update='momentum', sample_batches=False,
                                              num_epochs=100,
                                              verbose=False)

correct_loss = 0.494394
print('Final loss with momentum SGD: %f. We get: %f' % (loss_history[-1], correct_loss))

```

```

starting iteration 0
Final loss with momentum SGD: 0.494391. We get: 0.494394

```

The **RMSProp** update step is given as follows:

```

cache = decay_rate * cache + (1 - decay_rate) * dx**2
x += - learning_rate * dx / np.sqrt(cache + 1e-8)

```

Here, `decay_rate` is a hyperparameter and typical values are [0.9, 0.99, 0.999].

Implement the **RMSProp** update rule inside the train function and rerun the optimization:

```

In [8]: model = init_toy_model()
        trainer = ClassifierTrainer()

```

```

# call the trainer to optimize the loss
# Notice that we're using sample_batches=False, so we're performing Gradient Descent (no
best_model, loss_history, _, _ = trainer.train(X, y, X, y,
                                                model, two_layer_net,
                                                reg=0.001,
                                                learning_rate=1e-1, momentum=0.9, learning_
                                                update='rmsprop', sample_batches=False,
                                                num_epochs=100,
                                                verbose=False)

correct_loss = 0.439368
print('Final loss with RMSProp: %f. We get: %f' % (loss_history[-1], correct_loss))

starting iteration 0
Final loss with RMSProp: 0.439363. We get: 0.439368

```

## 6 Load the data

Now that you have implemented a two-layer network that passes gradient checks, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier.

```

In [9]: from cs231n.data_utils import load_CIFAR10

def get_CIFAR10_data(num_training=49000, num_validation=1000, num_test=1000):
    """
    Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
    it for the two-layer neural net classifier.
    """
    # Load the raw CIFAR-10 data
    cifar10_dir = 'cs231n/datasets/cifar-10-batches-py'
    X_train, y_train, X_test, y_test = load_CIFAR10(cifar10_dir)

    # Subsample the data
    mask = range(num_training, num_training + num_validation)
    X_val = X_train[mask]
    y_val = y_train[mask]
    mask = range(num_training)
    X_train = X_train[mask]
    y_train = y_train[mask]
    mask = range(num_test)
    X_test = X_test[mask]
    y_test = y_test[mask]

    # Normalize the data: subtract the mean image
    mean_image = np.mean(X_train, axis=0)
    X_train -= mean_image
    X_val -= mean_image
    X_test -= mean_image

```

```

    # Reshape data to rows
    X_train = X_train.reshape(num_training, -1)
    X_val = X_val.reshape(num_validation, -1)
    X_test = X_test.reshape(num_test, -1)

    return X_train, y_train, X_val, y_val, X_test, y_test

# Invoke the above function to get our data.
X_train, y_train, X_val, y_val, X_test, y_test = get_CIFAR10_data()
print('Train data shape: ', X_train.shape)
print('Train labels shape: ', y_train.shape)
print('Validation data shape: ', X_val.shape)
print('Validation labels shape: ', y_val.shape)
print('Test data shape: ', X_test.shape)
print('Test labels shape: ', y_test.shape)

```

```

Train data shape: (49000, 3072)
Train labels shape: (49000,)
Validation data shape: (1000, 3072)
Validation labels shape: (1000,)
Test data shape: (1000, 3072)
Test labels shape: (1000,)

```

## 7 Train a network

To train our network we will use SGD with momentum. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

```
In [10]: from cs231n.classifiers.neural_net import init_two_layer_model
```

```

model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train, X_val, y_val,
    model, two_layer_net,
    num_epochs=5, reg=1.0,
    momentum=0.9, learning_rate_decay = 0.95,
    learning_rate=1e-5, verbose=False)

```

```

starting iteration 0
starting iteration 100
starting iteration 200
starting iteration 300
starting iteration 400
starting iteration 500

```

```
starting iteration 600
starting iteration 700
starting iteration 800
starting iteration 900
starting iteration 1000
starting iteration 1100
starting iteration 1200
starting iteration 1300
starting iteration 1400
starting iteration 1500
starting iteration 1600
starting iteration 1700
starting iteration 1800
starting iteration 1900
starting iteration 2000
starting iteration 2100
starting iteration 2200
starting iteration 2300
starting iteration 2400
```

## 8 Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 0.37 on the validation set. This isn't very good.


One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

```
In [11]: # Plot the loss function and train / validation accuracies
plt.subplot(2, 1, 1)
plt.plot(loss_history)
plt.title('Loss history')
plt.xlabel('Iteration')
plt.ylabel('Loss')

plt.subplot(2, 1, 2)
plt.plot(train_acc)
plt.plot(val_acc)
plt.legend(['Training accuracy', 'Validation accuracy'], loc='lower right')
plt.xlabel('Epoch')
plt.ylabel('Clasification accuracy')

Out[11]: Text(0,0.5,'Clasification accuracy')
```



```
In [12]: from cs231n.vis_utils import visualize_grid

         # Visualize the weights of the network

def show_net_weights(model):
    plt.imshow(visualize_grid(model['W1'].T.reshape(-1, 32, 32, 3), padding=3).astype('float'))
    plt.gca().axis('off')
    plt.show()

show_net_weights(model)
```





## 9 Tune your hyperparameters

**What's wrong?.** Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

**Tuning.** Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might

also consider tuning the momentum and learning rate decay parameters, but you should be able to get good performance using the default values.

**Approximate results.** You should be aim to achieve a classification accuracy of greater than 50% on the validation set. Our best network gets over 56% on the validation set.

**Experiment:** Your goal in this exercise is to get as good of a result on CIFAR-10 as you can, with a fully-connected Neural Network. For every 1% above 56% on the Test set we will award you with one extra bonus point. Feel free to implement your own techniques (e.g. PCA to reduce dimensionality, or adding dropout, or adding features to the solver, etc.).

hyperparameters cell

```
1. Test = 41 num_epochs=10, reg=1, momentum=0.9, learning_rate_decay=0.95,
   learning_rate=1e-5, verbose=True)
```

```
2. 48.6
```

```
num_epochs=50, reg=1,
momentum=0.9,
learning_rate_decay=0.95,
learning_rate=1e-5, verbose=True)
```

```
In [19]: best_model = None # store the best model into this
```

```
#####
# TODO: Tune hyperparameters using the validation set. Store your best trained #
# model in best_model.                                                         #
#                                                                              #
# To help debug your network, it may help to use visualizations similar to the #
# ones we used above; these visualizations will have significant qualitative    #
# differences from the ones we saw above for the poorly tuned network.         #
#                                                                              #
# Tweaking hyperparameters by hand can be fun, but you might find it useful to #
# write code to sweep through possible combinations of hyperparameters        #
# automatically like we did on the previous assignment.                       #
#####
# input size, hidden size, number of classes
model = init_two_layer_model(32*32*3, 100, 10)
trainer = ClassifierTrainer()
best_model, loss_history, train_acc, val_acc = trainer.train(X_train, y_train,
                                                             X_val, y_val,
                                                             model, two_layer_net,
                                                             num_epochs=150, reg=1,
                                                             momentum=0.9,
                                                             learning_rate_decay=0.95,
                                                             learning_rate=1e-5, verbose=True)
#####
#                                     END OF YOUR CODE                         #
#####
```

starting iteration 0  
Finished epoch 0 / 150: cost 2.302591, train: 0.106000, val 0.071000, lr 1.000000e-05  
starting iteration 100  
starting iteration 200  
starting iteration 300  
starting iteration 400  
Finished epoch 1 / 150: cost 2.230374, train: 0.185000, val 0.179000, lr 9.500000e-06  
starting iteration 500  
starting iteration 600  
starting iteration 700  
starting iteration 800  
starting iteration 900  
Finished epoch 2 / 150: cost 2.061289, train: 0.235000, val 0.244000, lr 9.025000e-06  
starting iteration 1000  
starting iteration 1100  
starting iteration 1200  
starting iteration 1300  
starting iteration 1400  
Finished epoch 3 / 150: cost 1.927104, train: 0.291000, val 0.298000, lr 8.573750e-06  
starting iteration 1500  
starting iteration 1600  
starting iteration 1700  
starting iteration 1800  
starting iteration 1900  
Finished epoch 4 / 150: cost 1.875031, train: 0.339000, val 0.337000, lr 8.145063e-06  
starting iteration 2000  
starting iteration 2100  
starting iteration 2200  
starting iteration 2300  
starting iteration 2400  
Finished epoch 5 / 150: cost 1.707858, train: 0.324000, val 0.368000, lr 7.737809e-06  
starting iteration 2500  
starting iteration 2600  
starting iteration 2700  
starting iteration 2800  
starting iteration 2900  
Finished epoch 6 / 150: cost 1.747641, train: 0.364000, val 0.380000, lr 7.350919e-06  
starting iteration 3000  
starting iteration 3100  
starting iteration 3200  
starting iteration 3300  
starting iteration 3400  
Finished epoch 7 / 150: cost 1.854529, train: 0.420000, val 0.385000, lr 6.983373e-06  
starting iteration 3500  
starting iteration 3600  
starting iteration 3700  
starting iteration 3800  
starting iteration 3900

Finished epoch 8 / 150: cost 1.707482, train: 0.401000, val 0.405000, lr 6.634204e-06  
starting iteration 4000  
starting iteration 4100  
starting iteration 4200  
starting iteration 4300  
starting iteration 4400  
Finished epoch 9 / 150: cost 1.693510, train: 0.401000, val 0.412000, lr 6.302494e-06  
starting iteration 4500  
starting iteration 4600  
starting iteration 4700  
starting iteration 4800  
Finished epoch 10 / 150: cost 1.607555, train: 0.422000, val 0.433000, lr 5.987369e-06  
starting iteration 4900  
starting iteration 5000  
starting iteration 5100  
starting iteration 5200  
starting iteration 5300  
Finished epoch 11 / 150: cost 1.753437, train: 0.423000, val 0.437000, lr 5.688001e-06  
starting iteration 5400  
starting iteration 5500  
starting iteration 5600  
starting iteration 5700  
starting iteration 5800  
Finished epoch 12 / 150: cost 1.541589, train: 0.450000, val 0.446000, lr 5.403601e-06  
starting iteration 5900  
starting iteration 6000  
starting iteration 6100  
starting iteration 6200  
starting iteration 6300  
Finished epoch 13 / 150: cost 1.682181, train: 0.443000, val 0.445000, lr 5.133421e-06  
starting iteration 6400  
starting iteration 6500  
starting iteration 6600  
starting iteration 6700  
starting iteration 6800  
Finished epoch 14 / 150: cost 1.527999, train: 0.463000, val 0.457000, lr 4.876750e-06  
starting iteration 6900  
starting iteration 7000  
starting iteration 7100  
starting iteration 7200  
starting iteration 7300  
Finished epoch 15 / 150: cost 1.552910, train: 0.469000, val 0.455000, lr 4.632912e-06  
starting iteration 7400  
starting iteration 7500  
starting iteration 7600  
starting iteration 7700  
starting iteration 7800  
Finished epoch 16 / 150: cost 1.402946, train: 0.474000, val 0.463000, lr 4.401267e-06

starting iteration 7900  
starting iteration 8000  
starting iteration 8100  
starting iteration 8200  
starting iteration 8300  
Finished epoch 17 / 150: cost 1.644635, train: 0.470000, val 0.455000, lr 4.181203e-06  
starting iteration 8400  
starting iteration 8500  
starting iteration 8600  
starting iteration 8700  
starting iteration 8800  
Finished epoch 18 / 150: cost 1.615797, train: 0.452000, val 0.461000, lr 3.972143e-06  
starting iteration 8900  
starting iteration 9000  
starting iteration 9100  
starting iteration 9200  
starting iteration 9300  
Finished epoch 19 / 150: cost 1.479390, train: 0.460000, val 0.473000, lr 3.773536e-06  
starting iteration 9400  
starting iteration 9500  
starting iteration 9600  
starting iteration 9700  
Finished epoch 20 / 150: cost 1.526531, train: 0.497000, val 0.466000, lr 3.584859e-06  
starting iteration 9800  
starting iteration 9900  
starting iteration 10000  
starting iteration 10100  
starting iteration 10200  
Finished epoch 21 / 150: cost 1.515938, train: 0.464000, val 0.470000, lr 3.405616e-06  
starting iteration 10300  
starting iteration 10400  
starting iteration 10500  
starting iteration 10600  
starting iteration 10700  
Finished epoch 22 / 150: cost 1.551063, train: 0.471000, val 0.469000, lr 3.235335e-06  
starting iteration 10800  
starting iteration 10900  
starting iteration 11000  
starting iteration 11100  
starting iteration 11200  
Finished epoch 23 / 150: cost 1.622964, train: 0.494000, val 0.474000, lr 3.073569e-06  
starting iteration 11300  
starting iteration 11400  
starting iteration 11500  
starting iteration 11600  
starting iteration 11700  
Finished epoch 24 / 150: cost 1.496666, train: 0.475000, val 0.472000, lr 2.919890e-06  
starting iteration 11800

starting iteration 11900  
starting iteration 12000  
starting iteration 12100  
starting iteration 12200  
Finished epoch 25 / 150: cost 1.400345, train: 0.484000, val 0.474000, lr 2.773896e-06  
starting iteration 12300  
starting iteration 12400  
starting iteration 12500  
starting iteration 12600  
starting iteration 12700  
Finished epoch 26 / 150: cost 1.635315, train: 0.495000, val 0.470000, lr 2.635201e-06  
starting iteration 12800  
starting iteration 12900  
starting iteration 13000  
starting iteration 13100  
starting iteration 13200  
Finished epoch 27 / 150: cost 1.764031, train: 0.495000, val 0.471000, lr 2.503441e-06  
starting iteration 13300  
starting iteration 13400  
starting iteration 13500  
starting iteration 13600  
starting iteration 13700  
Finished epoch 28 / 150: cost 1.553671, train: 0.498000, val 0.473000, lr 2.378269e-06  
starting iteration 13800  
starting iteration 13900  
starting iteration 14000  
starting iteration 14100  
starting iteration 14200  
Finished epoch 29 / 150: cost 1.512731, train: 0.481000, val 0.486000, lr 2.259355e-06  
starting iteration 14300  
starting iteration 14400  
starting iteration 14500  
starting iteration 14600  
Finished epoch 30 / 150: cost 1.425067, train: 0.499000, val 0.477000, lr 2.146388e-06  
starting iteration 14700  
starting iteration 14800  
starting iteration 14900  
starting iteration 15000  
starting iteration 15100  
Finished epoch 31 / 150: cost 1.568940, train: 0.484000, val 0.481000, lr 2.039068e-06  
starting iteration 15200  
starting iteration 15300  
starting iteration 15400  
starting iteration 15500  
starting iteration 15600  
Finished epoch 32 / 150: cost 1.515593, train: 0.519000, val 0.480000, lr 1.937115e-06  
starting iteration 15700  
starting iteration 15800

starting iteration 15900  
starting iteration 16000  
starting iteration 16100  
Finished epoch 33 / 150: cost 1.590110, train: 0.484000, val 0.483000, lr 1.840259e-06  
starting iteration 16200  
starting iteration 16300  
starting iteration 16400  
starting iteration 16500  
starting iteration 16600  
Finished epoch 34 / 150: cost 1.564076, train: 0.505000, val 0.477000, lr 1.748246e-06  
starting iteration 16700  
starting iteration 16800  
starting iteration 16900  
starting iteration 17000  
starting iteration 17100  
Finished epoch 35 / 150: cost 1.421460, train: 0.492000, val 0.481000, lr 1.660834e-06  
starting iteration 17200  
starting iteration 17300  
starting iteration 17400  
starting iteration 17500  
starting iteration 17600  
Finished epoch 36 / 150: cost 1.548828, train: 0.505000, val 0.479000, lr 1.577792e-06  
starting iteration 17700  
starting iteration 17800  
starting iteration 17900  
starting iteration 18000  
starting iteration 18100  
Finished epoch 37 / 150: cost 1.515039, train: 0.498000, val 0.480000, lr 1.498903e-06  
starting iteration 18200  
starting iteration 18300  
starting iteration 18400  
starting iteration 18500  
starting iteration 18600  
Finished epoch 38 / 150: cost 1.535161, train: 0.518000, val 0.479000, lr 1.423957e-06  
starting iteration 18700  
starting iteration 18800  
starting iteration 18900  
starting iteration 19000  
starting iteration 19100  
Finished epoch 39 / 150: cost 1.450373, train: 0.515000, val 0.478000, lr 1.352760e-06  
starting iteration 19200  
starting iteration 19300  
starting iteration 19400  
starting iteration 19500  
Finished epoch 40 / 150: cost 1.762823, train: 0.508000, val 0.482000, lr 1.285122e-06  
starting iteration 19600  
starting iteration 19700  
starting iteration 19800

starting iteration 19900  
starting iteration 20000  
Finished epoch 41 / 150: cost 1.496068, train: 0.512000, val 0.485000, lr 1.220865e-06  
starting iteration 20100  
starting iteration 20200  
starting iteration 20300  
starting iteration 20400  
starting iteration 20500  
Finished epoch 42 / 150: cost 1.476640, train: 0.507000, val 0.484000, lr 1.159822e-06  
starting iteration 20600  
starting iteration 20700  
starting iteration 20800  
starting iteration 20900  
starting iteration 21000  
Finished epoch 43 / 150: cost 1.585615, train: 0.511000, val 0.487000, lr 1.101831e-06  
starting iteration 21100  
starting iteration 21200  
starting iteration 21300  
starting iteration 21400  
starting iteration 21500  
Finished epoch 44 / 150: cost 1.521527, train: 0.476000, val 0.483000, lr 1.046740e-06  
starting iteration 21600  
starting iteration 21700  
starting iteration 21800  
starting iteration 21900  
starting iteration 22000  
Finished epoch 45 / 150: cost 1.487731, train: 0.510000, val 0.485000, lr 9.944026e-07  
starting iteration 22100  
starting iteration 22200  
starting iteration 22300  
starting iteration 22400  
starting iteration 22500  
Finished epoch 46 / 150: cost 1.595450, train: 0.512000, val 0.481000, lr 9.446824e-07  
starting iteration 22600  
starting iteration 22700  
starting iteration 22800  
starting iteration 22900  
starting iteration 23000  
Finished epoch 47 / 150: cost 1.501601, train: 0.523000, val 0.485000, lr 8.974483e-07  
starting iteration 23100  
starting iteration 23200  
starting iteration 23300  
starting iteration 23400  
starting iteration 23500  
Finished epoch 48 / 150: cost 1.537054, train: 0.493000, val 0.485000, lr 8.525759e-07  
starting iteration 23600  
starting iteration 23700  
starting iteration 23800



starting iteration 23900  
starting iteration 24000  
Finished epoch 49 / 150: cost 1.642618, train: 0.511000, val 0.482000, lr 8.099471e-07  
starting iteration 24100  
starting iteration 24200  
starting iteration 24300  
starting iteration 24400  
Finished epoch 50 / 150: cost 1.540092, train: 0.510000, val 0.484000, lr 7.694498e-07  
starting iteration 24500  
starting iteration 24600  
starting iteration 24700  
starting iteration 24800  
starting iteration 24900  
Finished epoch 51 / 150: cost 1.589463, train: 0.516000, val 0.482000, lr 7.309773e-07  
starting iteration 25000  
starting iteration 25100  
starting iteration 25200  
starting iteration 25300  
starting iteration 25400  
Finished epoch 52 / 150: cost 1.565420, train: 0.503000, val 0.475000, lr 6.944284e-07  
starting iteration 25500  
starting iteration 25600  
starting iteration 25700  
starting iteration 25800  
starting iteration 25900  
Finished epoch 53 / 150: cost 1.611238, train: 0.527000, val 0.479000, lr 6.597070e-07  
starting iteration 26000  
starting iteration 26100  
starting iteration 26200  
starting iteration 26300  
starting iteration 26400  
Finished epoch 54 / 150: cost 1.528771, train: 0.520000, val 0.484000, lr 6.267216e-07  
starting iteration 26500  
starting iteration 26600  
starting iteration 26700  
starting iteration 26800  
starting iteration 26900  
Finished epoch 55 / 150: cost 1.509801, train: 0.465000, val 0.481000, lr 5.953856e-07  
starting iteration 27000  
starting iteration 27100  
starting iteration 27200  
starting iteration 27300  
starting iteration 27400  
Finished epoch 56 / 150: cost 1.498746, train: 0.515000, val 0.481000, lr 5.656163e-07  
starting iteration 27500  
starting iteration 27600  
starting iteration 27700  
starting iteration 27800

starting iteration 27900  
Finished epoch 57 / 150: cost 1.455091, train: 0.491000, val 0.484000, lr 5.373355e-07  
starting iteration 28000  
starting iteration 28100  
starting iteration 28200  
starting iteration 28300  
starting iteration 28400  
Finished epoch 58 / 150: cost 1.534273, train: 0.503000, val 0.490000, lr 5.104687e-07  
starting iteration 28500  
starting iteration 28600  
starting iteration 28700  
starting iteration 28800  
starting iteration 28900  
Finished epoch 59 / 150: cost 1.340238, train: 0.494000, val 0.489000, lr 4.849453e-07  
starting iteration 29000  
starting iteration 29100  
starting iteration 29200  
starting iteration 29300  
Finished epoch 60 / 150: cost 1.375773, train: 0.502000, val 0.486000, lr 4.606980e-07  
starting iteration 29400  
starting iteration 29500  
starting iteration 29600  
starting iteration 29700  
starting iteration 29800  
Finished epoch 61 / 150: cost 1.585073, train: 0.527000, val 0.485000, lr 4.376631e-07  
starting iteration 29900  
starting iteration 30000  
starting iteration 30100  
starting iteration 30200  
starting iteration 30300  
Finished epoch 62 / 150: cost 1.567944, train: 0.516000, val 0.487000, lr 4.157799e-07  
starting iteration 30400  
starting iteration 30500  
starting iteration 30600  
starting iteration 30700  
starting iteration 30800  
Finished epoch 63 / 150: cost 1.570394, train: 0.514000, val 0.486000, lr 3.949909e-07  
starting iteration 30900  
starting iteration 31000  
starting iteration 31100  
starting iteration 31200  
starting iteration 31300  
Finished epoch 64 / 150: cost 1.522651, train: 0.536000, val 0.483000, lr 3.752414e-07  
starting iteration 31400  
starting iteration 31500  
starting iteration 31600  
starting iteration 31700  
starting iteration 31800

Finished epoch 65 / 150: cost 1.529298, train: 0.493000, val 0.483000, lr 3.564793e-07  
starting iteration 31900  
starting iteration 32000  
starting iteration 32100  
starting iteration 32200  
starting iteration 32300  
Finished epoch 66 / 150: cost 1.537794, train: 0.488000, val 0.483000, lr 3.386554e-07  
starting iteration 32400  
starting iteration 32500  
starting iteration 32600  
starting iteration 32700  
starting iteration 32800  
Finished epoch 67 / 150: cost 1.518446, train: 0.526000, val 0.489000, lr 3.217226e-07  
starting iteration 32900  
starting iteration 33000  
starting iteration 33100  
starting iteration 33200  
starting iteration 33300  
Finished epoch 68 / 150: cost 1.491950, train: 0.533000, val 0.484000, lr 3.056365e-07  
starting iteration 33400  
starting iteration 33500  
starting iteration 33600  
starting iteration 33700  
starting iteration 33800  
Finished epoch 69 / 150: cost 1.419251, train: 0.529000, val 0.482000, lr 2.903546e-07  
starting iteration 33900  
starting iteration 34000  
starting iteration 34100  
starting iteration 34200  
Finished epoch 70 / 150: cost 1.539193, train: 0.519000, val 0.484000, lr 2.758369e-07  
starting iteration 34300  
starting iteration 34400  
starting iteration 34500  
starting iteration 34600  
starting iteration 34700  
Finished epoch 71 / 150: cost 1.519022, train: 0.487000, val 0.487000, lr 2.620451e-07  
starting iteration 34800  
starting iteration 34900  
starting iteration 35000  
starting iteration 35100  
starting iteration 35200  
Finished epoch 72 / 150: cost 1.305177, train: 0.515000, val 0.487000, lr 2.489428e-07  
starting iteration 35300  
starting iteration 35400  
starting iteration 35500  
starting iteration 35600  
starting iteration 35700  
Finished epoch 73 / 150: cost 1.615088, train: 0.515000, val 0.485000, lr 2.364957e-07

starting iteration 35800  
starting iteration 35900  
starting iteration 36000  
starting iteration 36100  
starting iteration 36200  
Finished epoch 74 / 150: cost 1.663850, train: 0.519000, val 0.486000, lr 2.246709e-07  
starting iteration 36300  
starting iteration 36400  
starting iteration 36500  
starting iteration 36600  
starting iteration 36700  
Finished epoch 75 / 150: cost 1.704994, train: 0.494000, val 0.490000, lr 2.134373e-07  
starting iteration 36800  
starting iteration 36900  
starting iteration 37000  
starting iteration 37100  
starting iteration 37200  
Finished epoch 76 / 150: cost 1.420432, train: 0.500000, val 0.482000, lr 2.027655e-07  
starting iteration 37300  
starting iteration 37400  
starting iteration 37500  
starting iteration 37600  
starting iteration 37700  
Finished epoch 77 / 150: cost 1.528244, train: 0.496000, val 0.489000, lr 1.926272e-07  
starting iteration 37800  
starting iteration 37900  
starting iteration 38000  
starting iteration 38100  
starting iteration 38200  
Finished epoch 78 / 150: cost 1.698223, train: 0.518000, val 0.491000, lr 1.829958e-07  
starting iteration 38300  
starting iteration 38400  
starting iteration 38500  
starting iteration 38600  
starting iteration 38700  
Finished epoch 79 / 150: cost 1.428640, train: 0.508000, val 0.488000, lr 1.738460e-07  
starting iteration 38800  
starting iteration 38900  
starting iteration 39000  
starting iteration 39100  
Finished epoch 80 / 150: cost 1.472300, train: 0.515000, val 0.490000, lr 1.651537e-07  
starting iteration 39200  
starting iteration 39300  
starting iteration 39400  
starting iteration 39500  
starting iteration 39600  
Finished epoch 81 / 150: cost 1.503550, train: 0.511000, val 0.485000, lr 1.568961e-07  
starting iteration 39700

starting iteration 39800  
starting iteration 39900  
starting iteration 40000  
starting iteration 40100  
Finished epoch 82 / 150: cost 1.530307, train: 0.509000, val 0.487000, lr 1.490513e-07  
starting iteration 40200  
starting iteration 40300  
starting iteration 40400  
starting iteration 40500  
starting iteration 40600  
Finished epoch 83 / 150: cost 1.551121, train: 0.528000, val 0.485000, lr 1.415987e-07  
starting iteration 40700  
starting iteration 40800  
starting iteration 40900  
starting iteration 41000  
starting iteration 41100  
Finished epoch 84 / 150: cost 1.486826, train: 0.534000, val 0.489000, lr 1.345188e-07  
starting iteration 41200  
starting iteration 41300  
starting iteration 41400  
starting iteration 41500  
starting iteration 41600  
Finished epoch 85 / 150: cost 1.417320, train: 0.499000, val 0.488000, lr 1.277928e-07  
starting iteration 41700  
starting iteration 41800  
starting iteration 41900  
starting iteration 42000  
starting iteration 42100  
Finished epoch 86 / 150: cost 1.563426, train: 0.521000, val 0.491000, lr 1.214032e-07  
starting iteration 42200  
starting iteration 42300  
starting iteration 42400  
starting iteration 42500  
starting iteration 42600  
Finished epoch 87 / 150: cost 1.463056, train: 0.499000, val 0.489000, lr 1.153330e-07  
starting iteration 42700  
starting iteration 42800  
starting iteration 42900  
starting iteration 43000  
starting iteration 43100  
Finished epoch 88 / 150: cost 1.545913, train: 0.530000, val 0.490000, lr 1.095664e-07  
starting iteration 43200  
starting iteration 43300  
starting iteration 43400  
starting iteration 43500  
starting iteration 43600  
Finished epoch 89 / 150: cost 1.352506, train: 0.508000, val 0.490000, lr 1.040880e-07  
starting iteration 43700

starting iteration 43800  
starting iteration 43900  
starting iteration 44000  
Finished epoch 90 / 150: cost 1.407770, train: 0.499000, val 0.490000, lr 9.888365e-08  
starting iteration 44100  
starting iteration 44200  
starting iteration 44300  
starting iteration 44400  
starting iteration 44500  
Finished epoch 91 / 150: cost 1.454014, train: 0.533000, val 0.490000, lr 9.393946e-08  
starting iteration 44600  
starting iteration 44700  
starting iteration 44800  
starting iteration 44900  
starting iteration 45000  
Finished epoch 92 / 150: cost 1.585857, train: 0.505000, val 0.486000, lr 8.924249e-08  
starting iteration 45100  
starting iteration 45200  
starting iteration 45300  
starting iteration 45400  
starting iteration 45500  
Finished epoch 93 / 150: cost 1.725863, train: 0.521000, val 0.489000, lr 8.478037e-08  
starting iteration 45600  
starting iteration 45700  
starting iteration 45800  
starting iteration 45900  
starting iteration 46000  
Finished epoch 94 / 150: cost 1.421567, train: 0.494000, val 0.489000, lr 8.054135e-08  
starting iteration 46100  
starting iteration 46200  
starting iteration 46300  
starting iteration 46400  
starting iteration 46500  
Finished epoch 95 / 150: cost 1.503165, train: 0.535000, val 0.490000, lr 7.651428e-08  
starting iteration 46600  
starting iteration 46700  
starting iteration 46800  
starting iteration 46900  
starting iteration 47000  
Finished epoch 96 / 150: cost 1.544809, train: 0.501000, val 0.490000, lr 7.268857e-08  
starting iteration 47100  
starting iteration 47200  
starting iteration 47300  
starting iteration 47400  
starting iteration 47500  
Finished epoch 97 / 150: cost 1.452518, train: 0.536000, val 0.490000, lr 6.905414e-08  
starting iteration 47600  
starting iteration 47700

starting iteration 47800  
starting iteration 47900  
starting iteration 48000  
Finished epoch 98 / 150: cost 1.446358, train: 0.527000, val 0.489000, lr 6.560143e-08  
starting iteration 48100  
starting iteration 48200  
starting iteration 48300  
starting iteration 48400  
starting iteration 48500  
Finished epoch 99 / 150: cost 1.557571, train: 0.504000, val 0.490000, lr 6.232136e-08  
starting iteration 48600  
starting iteration 48700  
starting iteration 48800  
starting iteration 48900  
Finished epoch 100 / 150: cost 1.393283, train: 0.522000, val 0.489000, lr 5.920529e-08  
starting iteration 49000  
starting iteration 49100  
starting iteration 49200  
starting iteration 49300  
starting iteration 49400  
Finished epoch 101 / 150: cost 1.591829, train: 0.526000, val 0.490000, lr 5.624503e-08  
starting iteration 49500  
starting iteration 49600  
starting iteration 49700  
starting iteration 49800  
starting iteration 49900  
Finished epoch 102 / 150: cost 1.595473, train: 0.522000, val 0.488000, lr 5.343278e-08  
starting iteration 50000  
starting iteration 50100  
starting iteration 50200  
starting iteration 50300  
starting iteration 50400  
Finished epoch 103 / 150: cost 1.347366, train: 0.538000, val 0.490000, lr 5.076114e-08  
starting iteration 50500  
starting iteration 50600  
starting iteration 50700  
starting iteration 50800  
starting iteration 50900  
Finished epoch 104 / 150: cost 1.470135, train: 0.509000, val 0.493000, lr 4.822308e-08  
starting iteration 51000  
starting iteration 51100  
starting iteration 51200  
starting iteration 51300  
starting iteration 51400  
Finished epoch 105 / 150: cost 1.557893, train: 0.516000, val 0.490000, lr 4.581193e-08  
starting iteration 51500  
starting iteration 51600  
starting iteration 51700

starting iteration 51800  
starting iteration 51900  
Finished epoch 106 / 150: cost 1.570598, train: 0.524000, val 0.490000, lr 4.352133e-08  
starting iteration 52000  
starting iteration 52100  
starting iteration 52200  
starting iteration 52300  
starting iteration 52400  
Finished epoch 107 / 150: cost 1.507779, train: 0.527000, val 0.490000, lr 4.134526e-08  
starting iteration 52500  
starting iteration 52600  
starting iteration 52700  
starting iteration 52800  
starting iteration 52900  
Finished epoch 108 / 150: cost 1.537996, train: 0.498000, val 0.492000, lr 3.927800e-08  
starting iteration 53000  
starting iteration 53100  
starting iteration 53200  
starting iteration 53300  
starting iteration 53400  
Finished epoch 109 / 150: cost 1.547336, train: 0.544000, val 0.491000, lr 3.731410e-08  
starting iteration 53500  
starting iteration 53600  
starting iteration 53700  
starting iteration 53800  
Finished epoch 110 / 150: cost 1.460860, train: 0.528000, val 0.491000, lr 3.544840e-08  
starting iteration 53900  
starting iteration 54000  
starting iteration 54100  
starting iteration 54200  
starting iteration 54300  
Finished epoch 111 / 150: cost 1.452390, train: 0.516000, val 0.490000, lr 3.367598e-08  
starting iteration 54400  
starting iteration 54500  
starting iteration 54600  
starting iteration 54700  
starting iteration 54800  
Finished epoch 112 / 150: cost 1.474742, train: 0.529000, val 0.491000, lr 3.199218e-08  
starting iteration 54900  
starting iteration 55000  
starting iteration 55100  
starting iteration 55200  
starting iteration 55300  
Finished epoch 113 / 150: cost 1.459167, train: 0.529000, val 0.488000, lr 3.039257e-08  
starting iteration 55400  
starting iteration 55500  
starting iteration 55600  
starting iteration 55700



starting iteration 55800  
Finished epoch 114 / 150: cost 1.602999, train: 0.519000, val 0.490000, lr 2.887294e-08  
starting iteration 55900  
starting iteration 56000  
starting iteration 56100  
starting iteration 56200  
starting iteration 56300  
Finished epoch 115 / 150: cost 1.390391, train: 0.521000, val 0.489000, lr 2.742929e-08  
starting iteration 56400  
starting iteration 56500  
starting iteration 56600  
starting iteration 56700  
starting iteration 56800  
Finished epoch 116 / 150: cost 1.402484, train: 0.521000, val 0.490000, lr 2.605783e-08  
starting iteration 56900  
starting iteration 57000  
starting iteration 57100  
starting iteration 57200  
starting iteration 57300  
Finished epoch 117 / 150: cost 1.362262, train: 0.509000, val 0.489000, lr 2.475494e-08  
starting iteration 57400  
starting iteration 57500  
starting iteration 57600  
starting iteration 57700  
starting iteration 57800  
Finished epoch 118 / 150: cost 1.350107, train: 0.535000, val 0.488000, lr 2.351719e-08  
starting iteration 57900  
starting iteration 58000  
starting iteration 58100  
starting iteration 58200  
starting iteration 58300  
Finished epoch 119 / 150: cost 1.396687, train: 0.527000, val 0.489000, lr 2.234133e-08  
starting iteration 58400  
starting iteration 58500  
starting iteration 58600  
starting iteration 58700  
Finished epoch 120 / 150: cost 1.580804, train: 0.515000, val 0.488000, lr 2.122426e-08  
starting iteration 58800  
starting iteration 58900  
starting iteration 59000  
starting iteration 59100  
starting iteration 59200  
Finished epoch 121 / 150: cost 1.595670, train: 0.514000, val 0.489000, lr 2.016305e-08  
starting iteration 59300  
starting iteration 59400  
starting iteration 59500  
starting iteration 59600  
starting iteration 59700

Finished epoch 122 / 150: cost 1.319490, train: 0.530000, val 0.488000, lr 1.915490e-08  
starting iteration 59800  
starting iteration 59900  
starting iteration 60000  
starting iteration 60100  
starting iteration 60200  
Finished epoch 123 / 150: cost 1.475045, train: 0.528000, val 0.488000, lr 1.819715e-08  
starting iteration 60300  
starting iteration 60400  
starting iteration 60500  
starting iteration 60600  
starting iteration 60700  
Finished epoch 124 / 150: cost 1.464098, train: 0.525000, val 0.489000, lr 1.728730e-08  
starting iteration 60800  
starting iteration 60900  
starting iteration 61000  
starting iteration 61100  
starting iteration 61200  
Finished epoch 125 / 150: cost 1.481332, train: 0.489000, val 0.490000, lr 1.642293e-08  
starting iteration 61300  
starting iteration 61400  
starting iteration 61500  
starting iteration 61600  
starting iteration 61700  
Finished epoch 126 / 150: cost 1.543631, train: 0.510000, val 0.490000, lr 1.560178e-08  
starting iteration 61800  
starting iteration 61900  
starting iteration 62000  
starting iteration 62100  
starting iteration 62200  
Finished epoch 127 / 150: cost 1.681906, train: 0.493000, val 0.490000, lr 1.482169e-08  
starting iteration 62300  
starting iteration 62400  
starting iteration 62500  
starting iteration 62600  
starting iteration 62700  
Finished epoch 128 / 150: cost 1.605404, train: 0.487000, val 0.489000, lr 1.408061e-08  
starting iteration 62800  
starting iteration 62900  
starting iteration 63000  
starting iteration 63100  
starting iteration 63200  
Finished epoch 129 / 150: cost 1.467696, train: 0.525000, val 0.490000, lr 1.337658e-08  
starting iteration 63300  
starting iteration 63400  
starting iteration 63500  
starting iteration 63600  
Finished epoch 130 / 150: cost 1.636468, train: 0.501000, val 0.490000, lr 1.270775e-08

starting iteration 63700  
starting iteration 63800  
starting iteration 63900  
starting iteration 64000  
starting iteration 64100  
Finished epoch 131 / 150: cost 1.606836, train: 0.527000, val 0.490000, lr 1.207236e-08  
starting iteration 64200  
starting iteration 64300  
starting iteration 64400  
starting iteration 64500  
starting iteration 64600  
Finished epoch 132 / 150: cost 1.767178, train: 0.521000, val 0.490000, lr 1.146875e-08  
starting iteration 64700  
starting iteration 64800  
starting iteration 64900  
starting iteration 65000  
starting iteration 65100  
Finished epoch 133 / 150: cost 1.414837, train: 0.526000, val 0.490000, lr 1.089531e-08  
starting iteration 65200  
starting iteration 65300  
starting iteration 65400  
starting iteration 65500  
starting iteration 65600  
Finished epoch 134 / 150: cost 1.518343, train: 0.517000, val 0.490000, lr 1.035054e-08  
starting iteration 65700  
starting iteration 65800  
starting iteration 65900  
starting iteration 66000  
starting iteration 66100  
Finished epoch 135 / 150: cost 1.560910, train: 0.499000, val 0.490000, lr 9.833015e-09  
starting iteration 66200  
starting iteration 66300  
starting iteration 66400  
starting iteration 66500  
starting iteration 66600  
Finished epoch 136 / 150: cost 1.508085, train: 0.500000, val 0.490000, lr 9.341365e-09  
starting iteration 66700  
starting iteration 66800  
starting iteration 66900  
starting iteration 67000  
starting iteration 67100  
Finished epoch 137 / 150: cost 1.569809, train: 0.531000, val 0.490000, lr 8.874296e-09  
starting iteration 67200  
starting iteration 67300  
starting iteration 67400  
starting iteration 67500  
starting iteration 67600  
Finished epoch 138 / 150: cost 1.610751, train: 0.523000, val 0.490000, lr 8.430581e-09

starting iteration 67700  
starting iteration 67800  
starting iteration 67900  
starting iteration 68000  
starting iteration 68100  
Finished epoch 139 / 150: cost 1.616450, train: 0.512000, val 0.489000, lr 8.009052e-09  
starting iteration 68200  
starting iteration 68300  
starting iteration 68400  
starting iteration 68500  
Finished epoch 140 / 150: cost 1.455042, train: 0.516000, val 0.489000, lr 7.608600e-09  
starting iteration 68600  
starting iteration 68700  
starting iteration 68800  
starting iteration 68900  
starting iteration 69000  
Finished epoch 141 / 150: cost 1.546349, train: 0.512000, val 0.489000, lr 7.228170e-09  
starting iteration 69100  
starting iteration 69200  
starting iteration 69300  
starting iteration 69400  
starting iteration 69500  
Finished epoch 142 / 150: cost 1.482789, train: 0.510000, val 0.490000, lr 6.866761e-09  
starting iteration 69600  
starting iteration 69700  
starting iteration 69800  
starting iteration 69900  
starting iteration 70000  
Finished epoch 143 / 150: cost 1.531323, train: 0.515000, val 0.490000, lr 6.523423e-09  
starting iteration 70100  
starting iteration 70200  
starting iteration 70300  
starting iteration 70400  
starting iteration 70500  
Finished epoch 144 / 150: cost 1.538728, train: 0.510000, val 0.490000, lr 6.197252e-09  
starting iteration 70600  
starting iteration 70700  
starting iteration 70800  
starting iteration 70900  
starting iteration 71000  
Finished epoch 145 / 150: cost 1.484660, train: 0.510000, val 0.490000, lr 5.887389e-09  
starting iteration 71100  
starting iteration 71200  
starting iteration 71300  
starting iteration 71400  
starting iteration 71500  
Finished epoch 146 / 150: cost 1.434962, train: 0.543000, val 0.490000, lr 5.593020e-09  
starting iteration 71600

```
starting iteration 71700
starting iteration 71800
starting iteration 71900
starting iteration 72000
Finished epoch 147 / 150: cost 1.476097, train: 0.498000, val 0.490000, lr 5.313369e-09
starting iteration 72100
starting iteration 72200
starting iteration 72300
starting iteration 72400
starting iteration 72500
Finished epoch 148 / 150: cost 1.551739, train: 0.536000, val 0.490000, lr 5.047701e-09
starting iteration 72600
starting iteration 72700
starting iteration 72800
starting iteration 72900
starting iteration 73000
Finished epoch 149 / 150: cost 1.474683, train: 0.500000, val 0.490000, lr 4.795316e-09
starting iteration 73100
starting iteration 73200
starting iteration 73300
starting iteration 73400
Finished epoch 150 / 150: cost 1.474601, train: 0.506000, val 0.490000, lr 4.555550e-09
finished optimization. best validation accuracy: 0.493000
```

```
In [20]: # visualize the weights
         show_net_weights(best_model)
```



## 10 Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set.

```
In [21]: scores_test = two_layer_net(X_test, best_model)
         print('Test accuracy: ', np.mean(np.argmax(scores_test, axis=1) == y_test))
```

Test accuracy: 0.495